



Bundesministerium  
für Bildung  
und Forschung



## b2luigi introduction

(For those already familiar with luigi)

Joint meeting: Luigi-based Workflow Management with b2luigi/law

---

Michael Eliachevitch

([✉meliache@uni-bonn.de](mailto:meliache@uni-bonn.de), [🐙meliache](https://www.github.com/meliache), [🐦elimik31](https://twitter.com/elimik31), [inmeliache](https://www.linkedin.com/company/meliache))

Thursday 18 November

Physikalisches Institut — Rheinische Friedrich-Wilhelms-Universität Bonn

- full documentation at [b2luigi.readthedocs.io](https://b2luigi.readthedocs.io)
- lightweight helper package for `luigi`, that...

- full documentation at [b2luigi.readthedocs.io](https://b2luigi.readthedocs.io)
- lightweight helper package for `luigi`, that...
- brings `batch to (2) luigi`
  - use existing tasks, be able to switch between batch systems (LSF, HTCondor, local, ...)
  - schedule and monitor thousands of parallel jobs

- full documentation at [b2luigi.readthedocs.io](https://b2luigi.readthedocs.io)
- lightweight helper package for `luigi`, that...
- brings `batch to (2) luigi`
  - use existing tasks, be able to switch between batch systems (LSF, HTCondor, local, ...)
  - schedule and monitor thousands of parallel jobs
- helps with the `bread` 🍞 and `butter` 🧈 in `luigi`  
→ e.g. data management / output handling

- full documentation at [b2luigi.readthedocs.io](https://b2luigi.readthedocs.io)
- lightweight helper package for **luigi**, that...
- brings **batch to (2) luigi**
  - use existing tasks, be able to switch between batch systems (LSF, HTCondor, local, ...)
  - schedule and monitor thousands of parallel jobs
- helps with the **bread** 🍞 **and butter** 🧈 in luigi  
→ e.g. data management / output handling
- was originally developed for **Belle II**
  - collection of **basf2** helper tasks
  - grid-submission via wrapper for Belle II-specific tool **gbasf2**

- created in 2018 by Nils Braun at KIT for his Belle II tracking studies at the KEKCC LSF
- users at KIT and Belle II started contributing, e.g.
  - 2019: Max Welsch adds HTCondor batch (inspired by `law`)
  - 2020: I add LCG support for Belle II via soft `gbasf2` wrapper
  - ...
- 2021: Nils leaves for industry and I take over as main developer
- currently still in beta (`v0.7.4`), but has several active users within Belle II

## The team

### Main developer

Michael Eliachevitch ([meliache](#))

### Original author

Nils Braun ([nils-braun](#))

### Features, fixing, help and testing

- Felix Metzner ([FelixMetzner](#))
- Patrick Ecker ([eckerpatrick](#))
- Jochen Gemmler
- Maximilian Welsch ([welschma](#))
- Kilian Lieret ([klieret](#))
- Sviatoslav Bilokin ([bilokin](#))
- Phil Grace ([philiptgrace](#))
- Anselm Baur ([anselmbaur](#))
- Moritz Bauer ([sognetic](#))
- Artur Gottmann ([ArturAkh](#))

I imagine our users

- work with python anyway

I imagine our users

- work with python anyway

*"I write my steering file in python, I want to submit my jobs from python without having to learn some submit file syntax."*



I imagine our users

- work with python anyway  
*"I write my steering file in python, I want to submit my jobs from python without having to learn some submit file syntax."*
- comfortable with python classes  
(understands inheritance, properties, methods)

I imagine our users

- work with python anyway  
*"I write my steering file in python, I want to submit my jobs from python without having to learn some submit file syntax."*
- comfortable with python classes  
(understands inheritance, properties, methods)
- want to use bare `luigi` but wish for more convenient batch submission  
→ drop-in replacement: `import b2luigi as luigi`

I imagine our users

- work with python anyway  
*"I write my steering file in python, I want to submit my jobs from python without having to learn some submit file syntax."*
- comfortable with python classes  
(understands inheritance, properties, methods)
- want to use bare `luigi` but wish for more convenient batch submission  
→ drop-in replacement: `import b2luigi as luigi`
- use `b2luigi` as a `tool` to build their own workflow  
→ many users use it very differently

I imagine our users

- work with python anyway  
*"I write my steering file in python, I want to submit my jobs from python without having to learn some submit file syntax."*
- comfortable with python classes  
(understands inheritance, properties, methods)
- want to use bare `luigi` but wish for more convenient batch submission  
→ drop-in replacement: `import b2luigi as luigi`
- use `b2luigi` as a `tool` to build their own workflow  
→ many users use it very differently
- are potential contributors

- luigi/contrib already contains many batch scheduling and monitoring tasks (b2luigi was inspired by sge and lsf implementations)
- b2luigi solves some limitations of those systems, the improvements are:
  - *submit many parallel jobs*  
in other batch implementations each job needs a monitoring task, but number of processes is limited on many systems  
→ b2luigi requires only single process on submitting machine

- luigi/contrib already contains many batch scheduling and monitoring tasks (b2luigi was inspired by sge and lsf implementations)
- b2luigi solves some limitations of those systems, the improvements are:
  - [submit many parallel jobs](#)  
in other batch implementations each job needs a monitoring task, but number of processes is limited on many systems  
→ b2luigi requires only single process on submitting machine
  - [batch-submit a large set of luigi tasks](#)  
in other implementation you often need to implement a `work()` function instead of `run()` or define an external command to run

- luigi/contrib already contains many batch scheduling and monitoring tasks (b2luigi was inspired by sge and lsf implementations)
- b2luigi solves some limitations of those systems, the improvements are:
  - **submit many parallel jobs**  
in other batch implementations each job needs a monitoring task, but number of processes is limited on many systems  
→ b2luigi requires only single process on submitting machine
  - **batch-submit a large set of luigi tasks**  
in other implementation you often need to implement a `work()` function instead of `run()` or define an external command to run
  - **flexibility in choosing a batch systems**  
→ you can write your task first, test locally and then choose or change the batch you process them on

- write normal (b2)luigi task

```
import b2luigi

@b2luigi.requires(SomeOtherTask, foo_parameter="bar")
class MyTask(b2luigi.Task):
    cut_value = b2luigi.IntParameter():
    input_file = b2luigi.Parameter(hashed=True):

    def run(self): # ...
    def output(self): # ...
```



- write normal (b2)luigi task

```
import b2luigi

@b2luigi.requires(SomeOtherTask, foo_parameter="bar")
class MyTask(b2luigi.Task):
    cut_value = b2luigi.IntParameter():
    input_file = b2luigi.Parameter(hashed=True):

    def run(self): # ...
    def output(self): # ...
```

- use `b2luigi.process(tasks)` to run tasks locally or on a batch system

```
if __name__ == '__main__':
    b2luigi.set_setting("batch_system", "htcondor") # "local" for local submission
    tasks = [MyTask(cut_value=0.5, input_file=input_file) for input_file in input_files]
    b2luigi.process(tasks, batch=True)
```


- `b2luigi.process()` adds a CLI interface to your script:

```
$ python3 steering_file.py --help
usage: htcondor_example.py [-h] [--show-output] [--test] [--batch] [--batch-runner] [--dry-run]
↪ [--scheduler-host SCHEDULER_HOST] [--scheduler-port SCHEDULER_PORT] [--task-id TASK_ID]


optional arguments:
-h, --help                show this help message and exit
--show-output             Instead of running the tasks, show which output files will/are created.
--test                   Run the task list in test mode by printing the log directly to the screen
                        instead of storing it in a file.
--batch                  Instead of running locally, try to submit the tasks to the batch system.
--batch-runner            Expert option to mark this worker as a batch runner.
--dry-run                Do not run any task but set the return value to 0, if the tasks are complete.
--scheduler-host SCHEDULER_HOST
                        If given, use this host as a central scheduler instead of a local one.
--scheduler-port SCHEDULER_PORT
                        If given, use the port on this host as a central scheduler instead of a local
                        ↪ one.
--task-id TASK_ID        EXPERT.
```

- select **run-mode**: batch, test, dry-run or show-output
- connect to **central scheduler**

- provide additional options to b2luigi
- Can be set via
  1. via class attributes  
(e.g. static properties, luigi parameters or property functions)
  2. via

```
b2luigi.set_setting(key, value)
```
  3. **settings.json** configuration file
-  issue #31: *Support luigi's own config*

## Example

- provide additional options to b2luigi
- Can be set via
  1. via class attributes (e.g. static properties, luigi parameters or property functions)
  2. via `b2luigi.set_setting(key, value)`
  3. `settings.json` configuration file
-  issue #31: *Support luigi's own config*

```
class RecoTask(b2luigi.DispatchableTask):
    result_dir = "/path/to/results"
    batch_system = b2luigi.Parameter(significant=False)

    @property
    def htcondor_settings(self):
        return {"+requestRuntime": int(self.get_nevents() *
        ↪ 0.2)}

if __name__ == '__main__':
    b2luigi.set_setting("env_script", "./setup.sh")
    b2luigi.process(RecoTask(batch_system="htcondor"))
```


settings.json:

```
{"log_dir" : "/path/to/logs"}
```

- define interface `BatchProcess` (not as `luigi.Task`) with

- `start_job()`
- `kill_job()`
- `get_job_status() -> JobStatus`

```
class JobStatus(enum.Enum):  
    running = "running"  
    successful = "successful"  
    aborted = "aborted"  
    idle = "idle"
```

-  issue #2: *Include new batch systems.*  
Contributions welcome!

`b2luigi.Task` is a supercharged version of `luigi.Task`

- methods to help with **data management**

```
import b2luigi
import random

class MyNumberTask(b2luigi.Task):
    random_seed = b2luigi.IntParameter()

    def output(self):
        # ./random_seed=<seed>/output_file.txt
        yield self.add_to_output("output_file.txt")

    def run(self):
        random.seed(self.random_seed)
        random_number = random.random()
        # ./random_seed=<seed>/output_file.txt
        out_path = self.get_output_file_name(
            "output_file.txt"
        )
        with open(out_path, "w") as f:
            f.write(f"{random_number}\n")
```

`b2luigi.Task` is a supercharged version of `luigi.Task`

- methods to help with **data management**
  - helps you organize your outputs with the structure

```
<result_dir>/param1=foo/param2=bar/.../filename
```

```
import b2luigi
import random

class MyNumberTask(b2luigi.Task):
    random_seed = b2luigi.IntParameter()

    def output(self):
        # ./random_seed=<seed>/output_file.txt
        yield self.add_to_output("output_file.txt")

    def run(self):
        random.seed(self.random_seed)
        random_number = random.random()
        # ./random_seed=<seed>/output_file.txt
        out_path = self.get_output_file_name(
            "output_file.txt"
        )
        with open(out_path, "w") as f:
            f.write(f"{random_number}\n")
```

`b2luigi.Task` is a supercharged version of `luigi.Task`

- methods to help with **data management**

- helps you organize your outputs with the structure

```
<result_dir>/param1=foo/param2=bar/.../filename
```

- `self.add_to_output("<filename>")`
  - yield this `self.output()` generate a `luigi.LocalTarget` with a path under `result_dir` that encodes the values of the luigi parameter.

```
import b2luigi
import random

class MyNumberTask(b2luigi.Task):
    random_seed = b2luigi.IntParameter()

    def output(self):
        # ./random_seed=<seed>/output_file.txt
        yield self.add_to_output("output_file.txt")

    def run(self):
        random.seed(self.random_seed)
        random_number = random.random()
        # ./random_seed=<seed>/output_file.txt
        out_path = self.get_output_file_name(
            "output_file.txt"
        )
        with open(out_path, "w") as f:
            f.write(f"{random_number}\n")
```



`b2luigi.Task` is a supercharged version of `luigi.Task`

- methods to help with **data management**

- helps you organize your outputs with the structure

```
<result_dir>/param1=foo/param2=bar/.../filename
```

- `self.add_to_output("<filename>")`
  - yield this `self.output()` generate a `luigi.LocalTarget` with a path under `result_dir` that encodes the values of the luigi parameter.
- `self.get_output_file_name("<filename>")` :
  - use this in other methods like `run()` to get the generated file path.

```
import b2luigi
import random

class MyNumberTask(b2luigi.Task):
    random_seed = b2luigi.IntParameter()

    def output(self):
        # ./random_seed=<seed>/output_file.txt
        yield self.add_to_output("output_file.txt")

    def run(self):
        random.seed(self.random_seed)
        random_number = random.random()
        # ./random_seed=<seed>/output_file.txt
        out_path = self.get_output_file_name(
            "output_file.txt"
        )
        with open(out_path, "w") as f:
            f.write(f"{random_number}\n")
```

- **Problem:** Many HEP-tasks are steering-files that run `unsafe C++` framework code that can crash (`segfault`) and stop all tasks from running

- **Problem:** Many HEP-tasks are steering-files that run `unsafe C++` framework code that can crash (`segfault`) and stop all tasks from running
- **Solution:** *Dispatchable tasks* that emulate batch submission on local computer and runs it in separate execution path
  - decorate your `run()` method with `@b2luigi.dispatch`
  - or inherit from `b2luigi.DispatchableTask` and implement `process()` method instead of `run()` (this is what `Basf2Task` and its implementations do)

- **Problem:** Long-running tasks often create output before filling it.  
→ If task fails during execution, but the output is already created, it's mistakenly marked as *complete* ("thanksgiving bug")

- **Problem:** Long-running tasks often create output before filling it.  
→ If task fails during execution, but the output is already created, it's mistakenly marked as *complete* ("thanksgiving bug")
- **Solution:** Use `@b2luigi.on_temporary_files` decorator, e.g. for `run()` or `process()`  
→ Modifies `get_output_file_name()` to return temporary file location, to which output is written first. After the decorated function was successful, the output is moved to the actual output location.

- development happens on github, if you want something add an Issue or fork and PR
- unit tests encouraged
  - core-functionality well-covered
  - batch-systems not so (but still WIP)
- github actions for CI
  - pre-commit for style and static syntax checking
  - run unittests and calculate coverage
- see development documentation for a guide how to contribute

Thanks for listening.  
It's time for live action!

Except where otherwise noted, this work is licensed under  
<https://creativecommons.org/licenses/by-sa/4.0>





Backup